

We've made changes to our [Terms of Service](#) and [Privacy Policy](#). They take effect on September 1, 2020, and we encourage you to review them. By continuing to use our services, you agree to the new Terms of Service and acknowledge the Privacy Policy applies to you.

[Continue](#)

You have 2 free stories left this month. Sign up and get an extra one for free.

Gini Index vs Information Entropy

Everything you need to know about impurity measurements and information gain



Andrew Hershy

Jul 10, 2019 · 7 min read ★



Source

Introduction:

We've made changes to our [Terms of Service](#) and [Privacy Policy](#). They take effect on September 1, 2020, and we encourage you to review them. By continuing to use our services, you agree to the new Terms of Service and acknowledge the Privacy Policy applies to you.

Continue

learning, but I think it adds some valuable perspective to those wanting to understand the foundational mechanics of how decision trees work.

. . .

Table of Contents

- Background
- Gini Intuition
- Entropy Intuition
- Visualization
- Information Gain Comparison
- Practical Takeaways

. . .

Background:

Decision trees recursively split features with regard to their target variable's "purity". The entire algorithm is designed to optimize each split on maximizing purity... What is purity? Purity can be thought of as how homogenized the groupings are. You will see in some example below what I mean:

If we have 4 red gumballs and 0 blue gumballs, that group of 4 is 100% pure, based on color as the target.

If we have 2 red and 2 blue, that group is 100% impure.

If we have 3 red and 1 blue, that group is either 75% or 81% pure, if we use Gini or Entropy respectively.

We've made changes to our [Terms of Service](#) and [Privacy Policy](#). They take effect on September 1, 2020, and we encourage you to review them. By continuing to use our services, you agree to the new Terms of Service and acknowledge the Privacy Policy applies to you.

Continue

• • •

Gini Index Intuition:

$$Gini = 1 - \sum_j p_j^2$$

Let's start with Gini Index, as it's a bit easier to understand. According to Wikipedia, the goal is to “*measure how often a randomly chosen element from the set would be incorrectly labeled*”[1].

To visualize this, let's go back to the gumball examples. If we decided to arbitrarily label all 4 gumballs as red, how often would one of the gumballs be incorrectly labeled?

4 red and 0 blue:

$$Gini\ Index = 1 - (probability_red^2 + probability_blue^2) = 1 - (1^2 + 0^2) = 0$$

The impurity measurement is 0 because we would never incorrectly label any of the 4 red gumballs here. If we arbitrarily chose to label all the balls 'blue', then our index would still be 0, because we would always incorrectly label the gumballs.

The gini score is always the same no matter what arbitrary class you take the probabilities of because they always add to 0 in the formula above.

A gini score of 0 is the most pure score possible.

2 red and 2 blue:

$$Gini\ Index = 1 - (probability_red^2 + probability_blue^2) = 1 - (0.5^2 + 0.5^2) = 0.5$$

We've made changes to our [Terms of Service](#) and [Privacy Policy](#). They take effect on September 1, 2020, and we encourage you to review them. By continuing to use our services, you agree to the new Terms of Service and acknowledge the Privacy Policy applies to you.

Continue

Dividing gini scores by 0.5 can help intuitively understand what the score represents. $0.5/0.5 = 1$, meaning the grouping is as impure as possible (in a group with just 2 outcomes).

3 red and 1 blue:

$$\text{Gini Index} = 1 - (\text{probability_red}^2 + \text{probability_blue}^2) = 1 - (0.75^2 + 0.25^2) = 0.375$$

The impurity measurement here is 0.375. If we divide this by 0.5 for more intuitive understanding we will get 0.75, which is the probability of incorrectly/correctly labeling.

...

Entropy Intuition:

$$\text{Entropy} = - \sum_j p_j \log_2 p_j$$

Entropy is more computationally heavy due to the log in the equation. Like gini, The basic idea is to gauge the disorder of a grouping by the target variable. Instead of utilizing simple probabilities, this method takes the log base2 of the probabilities (you can use any log base, however, as long as you're consistent). The entropy equation uses logarithms because of many advantageous properties. The Main advantage is the additive property it provides. These MIT lecture notes will help grasp the concept more clearly (pg8) [2].

Let's visualize how entropy works with the same gumball scenarios:

4 red and 0 blue:

We've made changes to our [Terms of Service](#) and [Privacy Policy](#). They take effect on September 1, 2020, and we encourage you to review them. By continuing to use our services, you agree to the new Terms of Service and acknowledge the Privacy Policy applies to you.

Continue

purity score using information entropy.

2 red and 2 blue:

$$\text{Entropy} = [(\text{probability_red}) * \log_2(\text{probability_red})] - [(\text{probability_blue}) * \log_2(\text{probability of blue})] = [(2/4) * \log_2(2/4)] - [(2/4) * \log_2(2/4)] = 1$$

The impurity measurement is 1 here, as it's the maximum impurity obtainable.

3 red and 1 blue:

$$\text{Entropy} = [(\text{probability_red}) * \log_2(\text{probability_red})] - [(\text{probability_blue}) * \log_2(\text{probability of blue})] = [(3/4) * \log_2(3/4)] - [(1/4) * \log_2(1/4)] = 0.811$$

The purity/impurity measurement is 0.811 here, a bit worse than the gini score.

. . .

Visualization

Let's visualize both the Gini and Entropy curves with some code in python:

Gini:

Below we are making a function to automate gini calculations.

```
#Gini Function
#a and b are the quantities of each class
def gini(a,b):
    a1 = (a/(a+b))**2
    b1 = (b/(a+b))**2
    return 1 - (a1 + b1)
```

We've made changes to our [Terms of Service](#) and [Privacy Policy](#). They take effect on September 1, 2020, and we encourage you to review them. By continuing to use our services, you agree to the new Terms of Service and acknowledge the Privacy Policy applies to you.

Continue

```
#Blank lists
gini_list = []
blue_list = []
red_list = []
blue_prob_list = []

#Looping Gini function on random blue and red float amounts
for x in range (10000):
    blue = random.uniform(0, 4)
    red = abs(4-blue)
    a = gini(red,blue)
    b = blue/(blue+red)
    gini_list.append(a)
    blue_list.append(blue)
    red_list.append(red)
    blue_prob_list.append(b)

#Dataframe of amount of blue, red, Probability of blue, and gini
score
df = pd.DataFrame({"Blue": blue_list, "Red": red_list, "Gini Score":
gini_list, "Probability of Blue": blue_prob_list})
df = df[['Red', 'Blue', 'Probability of Blue', 'Gini Score']]
df
```

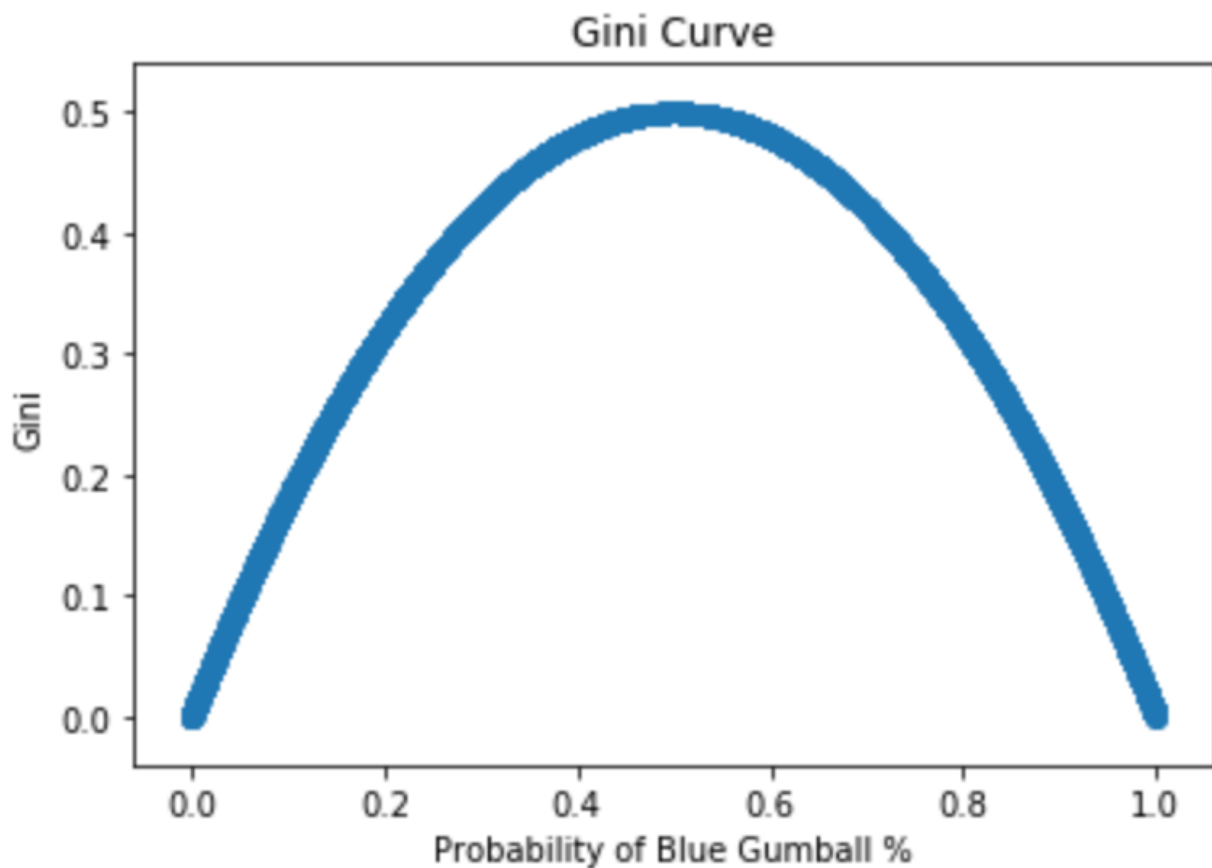
The beginning of the dataframe is below. The other 9,994 rows couldn't fit.

	Red	Blue	Probability of Blue	Gini Score
0	0.575254	3.424746	0.856186	0.246262
1	1.602532	2.397468	0.599367	0.480252
2	2.760943	1.239057	0.309764	0.427621
3	2.640753	1.359247	0.339812	0.448679
4	3.549481	0.450519	0.112630	0.199889
5	0.719211	3.280789	0.820197	0.294948

We've made changes to our [Terms of Service](#) and [Privacy Policy](#). They take effect on September 1, 2020, and we encourage you to review them. By continuing to use our services, you agree to the new Terms of Service and acknowledge the Privacy Policy applies to you.

Continue

```
plt.ylabel('Gini')
plt.title('Gini Curve')
```



Entropy:

```
#Gini Function
#a and b are the quantities of each class. Base is the log base
input.
def entropy(base,a,b):
    try:
        var = abs(((a)/(a+b)) * log(((a)/(a+b)),base)) -
        (((b)/(a+b)) * log(((b)/(a+b)),base))
        return var
    except (ValueError):
        return 0

#Blank lists
ent_list = []
blue_list = []
```

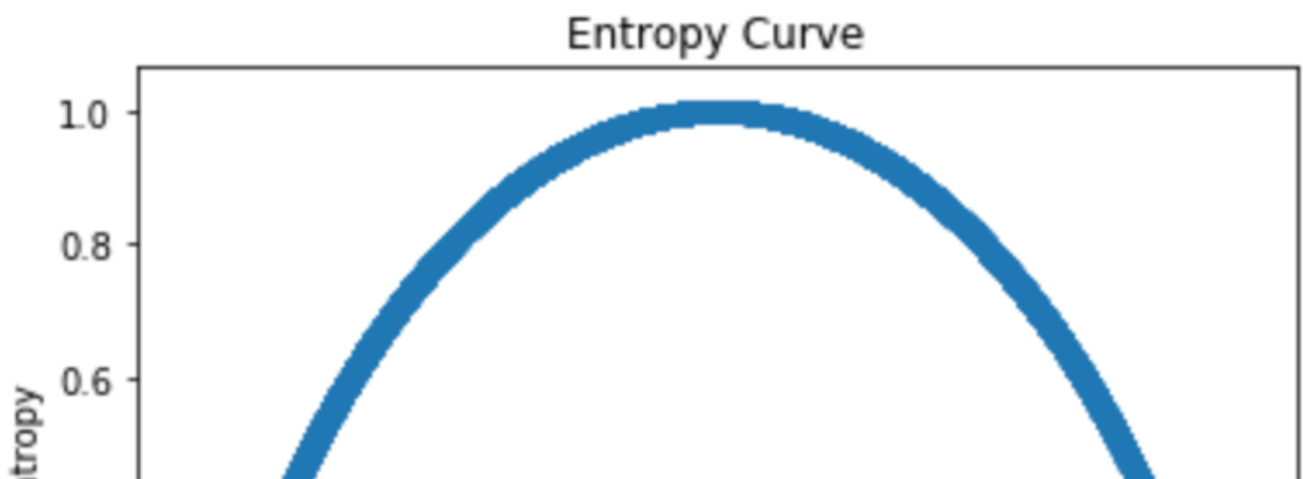
We've made changes to our [Terms of Service](#) and [Privacy Policy](#). They take effect on September 1, 2020, and we encourage you to review them. By continuing to use our services, you agree to the new Terms of Service and acknowledge the Privacy Policy applies to you.

Continue

```
red = abs(4-blue)
a = entropy(2, red, blue)
b = blue/(blue+red)
ent_list.append(a)
blue_list.append(blue)
red_list.append(red)
blue_prob_list.append(b)
```


```
df = pd.DataFrame({"Blue": blue_list, "Red": red_list, "Entropy":
ent_list, "Probability of Blue": blue_prob_list})
df = df[['Red', 'Blue', 'Probability of Blue', 'Entropy']]
df
```

	Red	Blue	Probability of Blue	Entropy
0	0.280530	3.719470	0.929867	0.366418
1	2.581253	1.418747	0.354687	0.938184
2	1.799219	2.200781	0.550195	0.992718
3	3.823594	0.176406	0.044101	0.260791
4	0.161097	3.838903	0.959726	0.243548
5	1.971193	2.028807	0.507202	0.999850



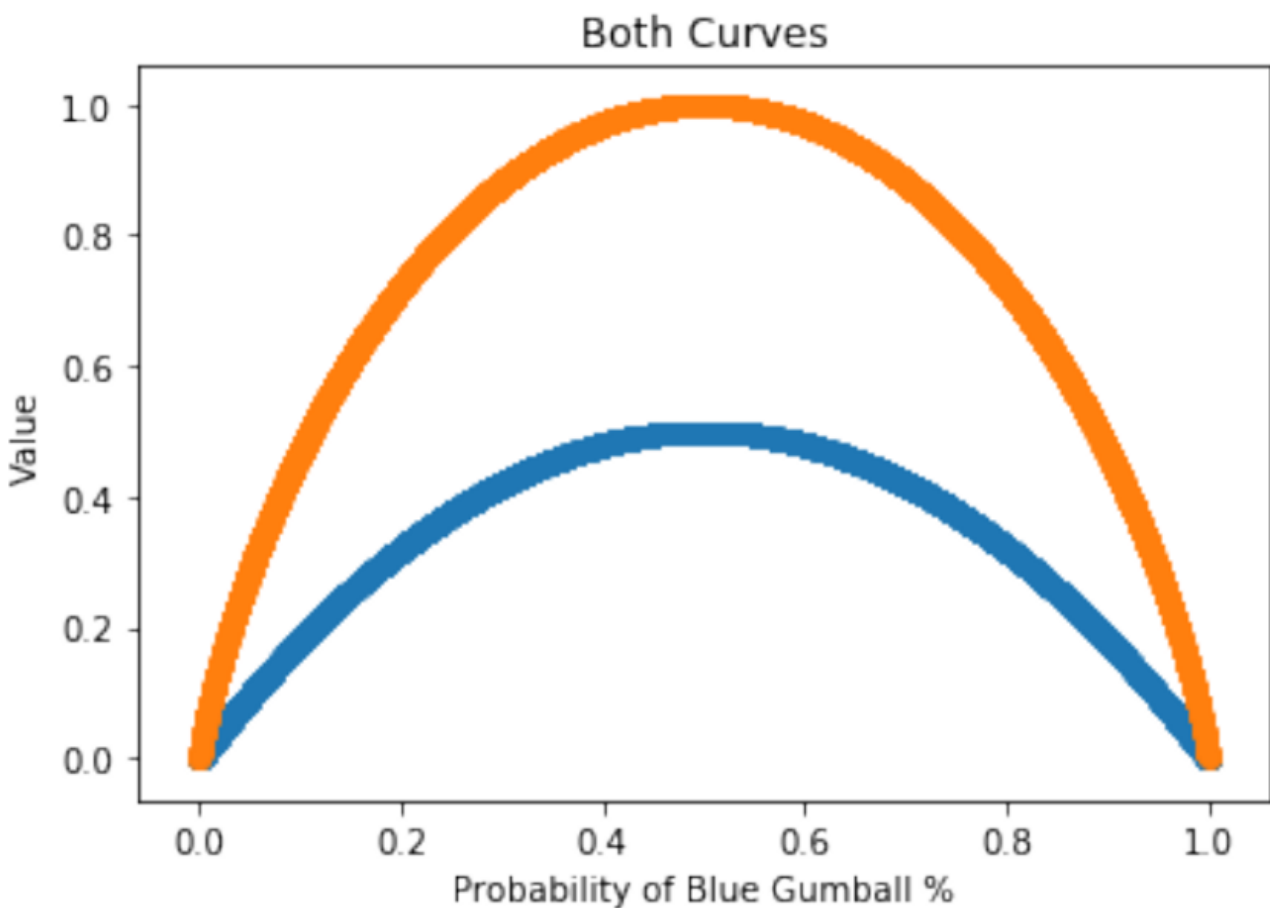
We've made changes to our [Terms of Service](#) and [Privacy Policy](#). They take effect on September 1, 2020, and we encourage you to review them. By continuing to use our services, you agree to the new Terms of Service and acknowledge the Privacy Policy applies to you.

Continue



Comparison:

Gini is blue, Entropy is orange. You will see how these differences are exemplified in information gain in the next section!



...

Information Gain

We've made changes to our [Terms of Service](#) and [Privacy Policy](#). They take effect on September 1, 2020, and we encourage you to review them. By continuing to use our services, you agree to the new Terms of Service and acknowledge the Privacy Policy applies to you.

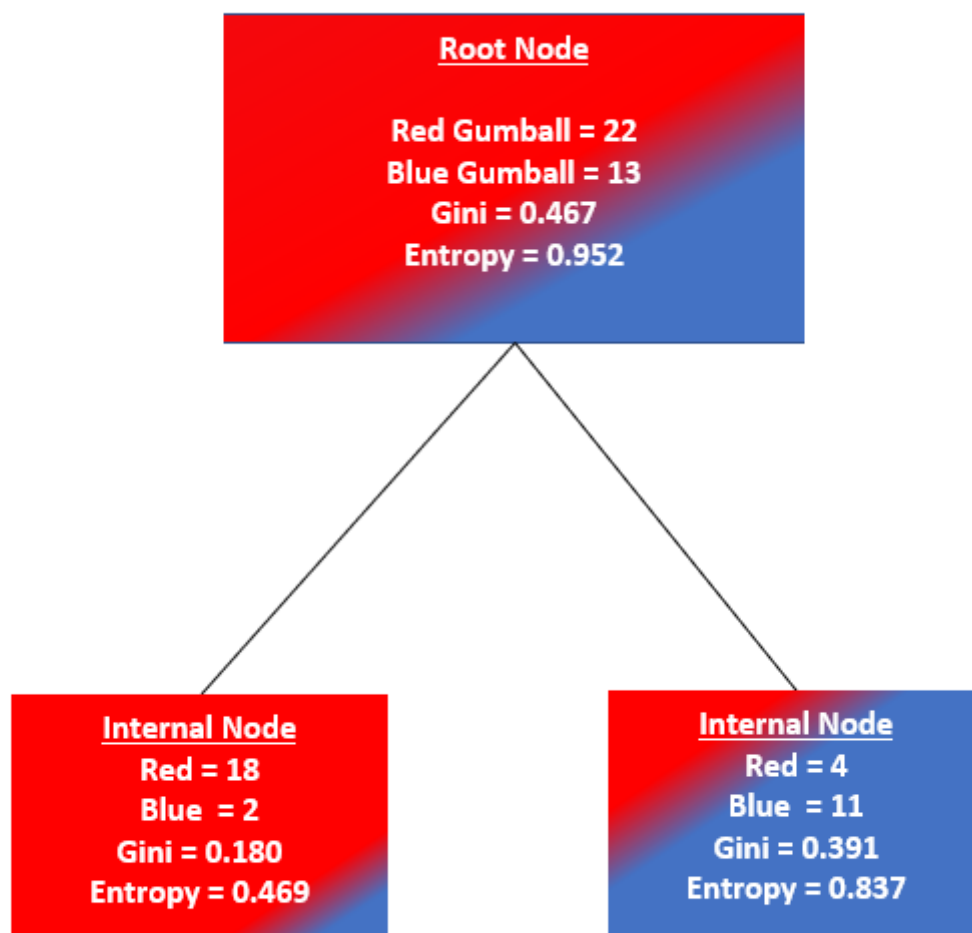
Continue

Information Gain (Parent, Child)

$$= Entropy(parent) - [p1(c1) * entropy(c1) + p(c2) * entropy(c2) + ...]$$

Source [3]

In the example below, we are splitting gumball preference by a particular attribute (let's say **size** of the gumball). It gives us the parent/child node relationship:



Gini Info Gain (from the equation above)

We've made changes to our [Terms of Service](#) and [Privacy Policy](#). They take effect on September 1, 2020, and we encourage you to review them. By continuing to use our services, you agree to the new Terms of Service and acknowledge the Privacy Policy applies to you.

Continue

#Running Function

```
gini_info_gain(22,13,18,2,4,11)
```

= 0.196 Gini Information gain

0.196/0.467 = 41.97% Gain

Entropy Info Gain (from the equation above)

#Defining Entropy info gain function:

```
def entropy_info_gain(base,pa,pb,c1a,c1b,c2a,c2b):
    return (entropy(base,pa,pb)) -
    (((c1a+c1b)/(pa+pb))*entropy(base,c1a,c1b)) +
    (((c2a+c2b)/(pa+pb))*entropy(base,c2a,c2b))
```

#Running Function

```
entropy_info_gain(2,22,13,18,2,4,11)
```

= 0.325 Entropy Information gain

0.325/0.952 = 34.14% Gain

Gini has a higher information gain measurement, for this example.

• • •

Final Takeaways:

1. Gini's maximum impurity is 0.5 and maximum purity is 0
2. Entropy's maximum impurity is 1 and maximum purity is 0
3. Different decision tree algorithms utilize different impurity metrics: CART uses Gini; ID3 and C4.5 use Entropy. This is worth looking into before you use decision

We've made changes to our [Terms of Service](#) and [Privacy Policy](#). They take effect on September 1, 2020, and we encourage you to review them. By continuing to use our services, you agree to the new Terms of Service and acknowledge the Privacy Policy applies to you.

Continue

• • •

Sources:

[1] https://en.wikipedia.org/wiki/Decision_tree_learning#Gini_impurity

[2] <http://web.mit.edu/6.02/www/f2011/handouts/2.pdf>

[3] Provost, F., & Fawcett, T. (2013). *Data Science for Business: What you need to know about data mining and data-analytic thinking*. Köln, CA: O`Reilly.

• • •

Please subscribe if you found this helpful. If you enjoy my content, please check out a few other projects:

Is Random Forest better than Logistic Regression? (a comparison)

Excel vs SQL: A Conceptual Comparison

Predicting Cancer with Logistic Regression in Python

Optimize your Investments using Math and Python

Calculating R-squared from scratch (using python)

Word Clouds in Python: Comprehensive Example

Sign up for The Daily Pick

By Towards Data Science

Hands-on real-world examples, research, tutorials, and cutting-edge techniques delivered Monday to Thursday. Make learning your daily ritual. [Take a look](#)

Get this newsletter

Create a free Medium account to get The Daily Pick in your inbox.

We've made changes to our [Terms of Service](#) and [Privacy Policy](#). They take effect on September 1, 2020, and we encourage you to review them. By continuing to use our services, you agree to the new Terms of Service and acknowledge the Privacy Policy applies to you.

[Continue](#)[About](#) [Help](#) [Legal](#)

Get the Medium app

